

Laplacian Eigenmaps for Dimensionality Reduction

Differential Methods for AI

Sergio Herreros Pérez, Mario Kroll Merino, Carlos Mazuecos Reíllo,
José Andrés Ridruejo Tuñón

Laplacian Eigenmaps: The Core Intuition

Goal

Map high-dimensional data $x_1, \dots, x_N \in \mathbb{R}^D$ to low-dimensional points $y_1, \dots, y_N \in \mathbb{R}^m$ (where $m \ll D$) while preserving **local neighborhood information**.

The Fundamental Idea:

- If two points x_i and x_j are close in the original space (high similarity), they should be mapped close together in the new space.
- We treat this as an optimization problem minimizing the stretching of edges in the neighborhood graph.

Step 1: Graph Construction & Weights

First, we define the geometry of the data manifold.

- ① **Adjacency (Neighbors):** Use k -Nearest Neighbors (k -NN). To ensure the graph is undirected, we enforce symmetry:

Connect i and j if $i \in kNN(j)$ **OR** $j \in kNN(i)$.

- ② **Weights (W):** Assign weights W_{ij} (Heat Kernel):

$$W_{ij} = e^{-\frac{\|x_i - x_j\|^2}{t}} \quad \text{if connected, else 0.}$$

- ③ **Degree Matrix (D):** Diagonal matrix measuring density.

$$D_{ii} = \sum_j W_{ij}$$

Step 2: The Optimization Problem (1D Case)

Let's try to map the data to a **single line** ($m = 1$). Let

$\mathbf{y} = (y_1, y_2, \dots, y_N)^T$ be the coordinate vector.

We want to minimize the weighted distance between neighbors:

Objective Function

$$\min_{\mathbf{y}} \sum_{i,j} (y_i - y_j)^2 W_{ij}$$

- If W_{ij} is large (points are close), y_i and y_j **must** be close to minimize the cost.
- If $W_{ij} = 0$, the distance $(y_i - y_j)$ doesn't matter.

Step 3: From Summation to Matrix Form

We can rewrite the sum using algebraic manipulation:

$$\begin{aligned}\sum_{i,j} (y_i - y_j)^2 W_{ij} &= \sum_{i,j} (y_i^2 + y_j^2 - 2y_i y_j) W_{ij} \\&= \sum_i y_i^2 D_{ii} + \sum_j y_j^2 D_{jj} - 2 \sum_{i,j} y_i y_j W_{ij} \\&= 2\mathbf{y}^T D \mathbf{y} - 2\mathbf{y}^T W \mathbf{y} \\&= 2\mathbf{y}^T (D - W)\mathbf{y}\end{aligned}$$

The Graph Laplacian

Since $L = D - W$, the objective function becomes:

$$\text{Cost}(\mathbf{y}) = 2\mathbf{y}^T L \mathbf{y}$$

Step 4: Constraints and Lagrange Multipliers

To prevent all points from collapsing to a single point, we impose a constraint to fix the scale and handle density:

$$\mathbf{y}^T D \mathbf{y} = 1$$

Now we solve using **Lagrange Multipliers**:

$$\mathcal{L}(\mathbf{y}, \lambda) = \mathbf{y}^T L \mathbf{y} - \lambda(\mathbf{y}^T D \mathbf{y} - 1)$$

Taking the derivative with respect to \mathbf{y} and setting to 0:

$$2L\mathbf{y} - 2\lambda D\mathbf{y} = 0 \implies \boxed{L\mathbf{y} = \lambda D\mathbf{y}}$$

This is the **Generalized Eigenvalue Problem**.

Step 5: Why Smallest Eigenvalues?

We have found that the optimal \mathbf{y} must be an eigenvector. But which one? Let's check the cost of a specific eigenvector solution \mathbf{y} :

$$\text{Cost} = \mathbf{y}^T L \mathbf{y}$$

Since $L\mathbf{y} = \lambda D\mathbf{y}$, we substitute:

$$\mathbf{y}^T L \mathbf{y} = \mathbf{y}^T (\lambda D \mathbf{y}) = \lambda (\mathbf{y}^T D \mathbf{y})$$

Using our constraint $\mathbf{y}^T D \mathbf{y} = 1$:

$$\text{Cost} = \lambda$$

Conclusion

To **minimize the cost**, we must choose the eigenvectors with the **smallest eigenvalues** λ .

Step 6: Generalizing to m Dimensions

- $\lambda_0 = 0$: Corresponds to eigenvector **1** (constant vector). Maps all points to a single spot. **Discard it.**
- $\lambda_1, \dots, \lambda_m$: The smallest non-zero eigenvalues.

The Embedding Map: For embedding into m dimensions, we use the eigenvectors $\mathbf{f}_1, \dots, \mathbf{f}_m$ associated with $0 < \lambda_1 \leq \dots \leq \lambda_m$.

$$x_i \mapsto (\mathbf{f}_1(i), \mathbf{f}_2(i), \dots, \mathbf{f}_m(i))$$

Summary

The coordinates in the low-dimensional space are literally the values of the Laplacian eigenvectors.

Laplace–Beltrami: From Continuous to Discrete

Continuous Manifold

Minimize Gradient Energy

$$\min_{\|u\|=1} \int_{\Omega} \|\nabla u\|^2 dx$$

↓ leads to

Laplace–Beltrami Operator

$$-\Delta u = \lambda u$$

Discrete Graph

Minimize Edge Tension

$$\min_{\mathbf{f}^T D \mathbf{f} = 1} \sum_{i,j} W_{ij} (f_i - f_j)^2$$

↓ leads to

Graph Laplacian

$$L\mathbf{f} = \lambda D\mathbf{f}$$

The Correspondence

Local Variation $(\nabla u)^2$	\longleftrightarrow	Edge Difference $(f_i - f_j)^2$
Laplace Operator $-\Delta$	\longleftrightarrow	Graph Matrix $L = D - W$
Manifold Harmonics	\longleftrightarrow	Graph Eigenvectors

Laplacian Eigenmaps in Graph-Based Learning

- Laplacian Eigenmaps provide the spectral foundation for many graph learning algorithms:
 - **Spectral Clustering:** uses the smallest non-zero Laplacian eigenvectors to partition graphs.
 - **Graph Convolutional Networks (GCNs):** graph convolutions are defined using the Laplacian spectrum.
- The Laplacian encodes the local geometry of the dataset by capturing similarity relations through weights w_{ij} .
- Eigenvectors provide smooth representations that respect this geometry.

Connection to the Finite Element Method (FEM)

Shared Variational Principle

Both FEM and Laplacian Eigenmaps are based on minimizing an energy associated with the Laplacian:

$$\int_{\Omega} |\nabla u|^2 \longleftrightarrow \sum_{i,j} w_{ij} (f_i - f_j)^2 = f^T L f.$$

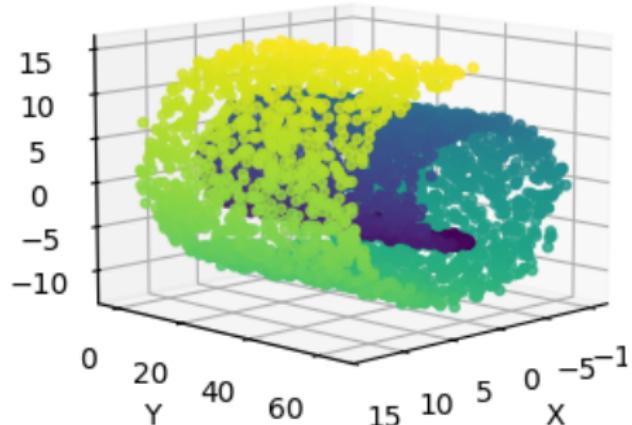
- FEM approximates Δ on a mesh; Laplacian Eigenmaps approximate Δ on a **similarity graph**.
- In both cases, low-energy eigenfunctions correspond to the smoothest modes.
- This provides a rigorous bridge between PDE-based models and graph learning.

From Continuous Geometry to Discrete Data

- Laplacian Eigenmaps allow transferring ideas from differential geometry to data analysis:
 - **Diffusion:** heat flow on a manifold \leftrightarrow diffusion processes on graphs.
 - **Smoothness:** low-oscillation eigenfunctions \leftrightarrow low-variation graph signals.
 - **Eigenmodes:** Laplace–Beltrami eigenfunctions \leftrightarrow graph Laplacian eigenvectors.
- This creates a unified framework connecting:
 - PDEs and variational principles,
 - manifold geometry,
 - and machine learning on graphs.
- The result: geometric structure of data becomes accessible even in discrete, high-dimensional settings.

Swiss Roll Dataset

Swiss Roll - View 1



Swiss Roll - View 2

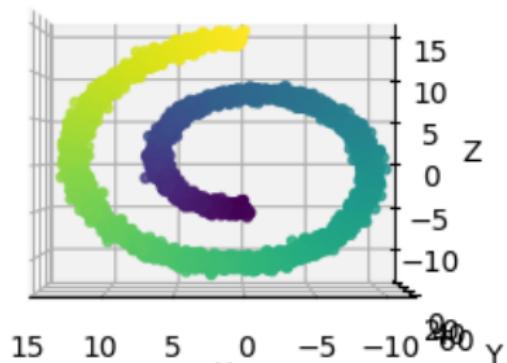


Figure: 3D Visualization of the Swiss Roll Dataset

Laplacian Eigenmaps Embedding of Swiss Roll

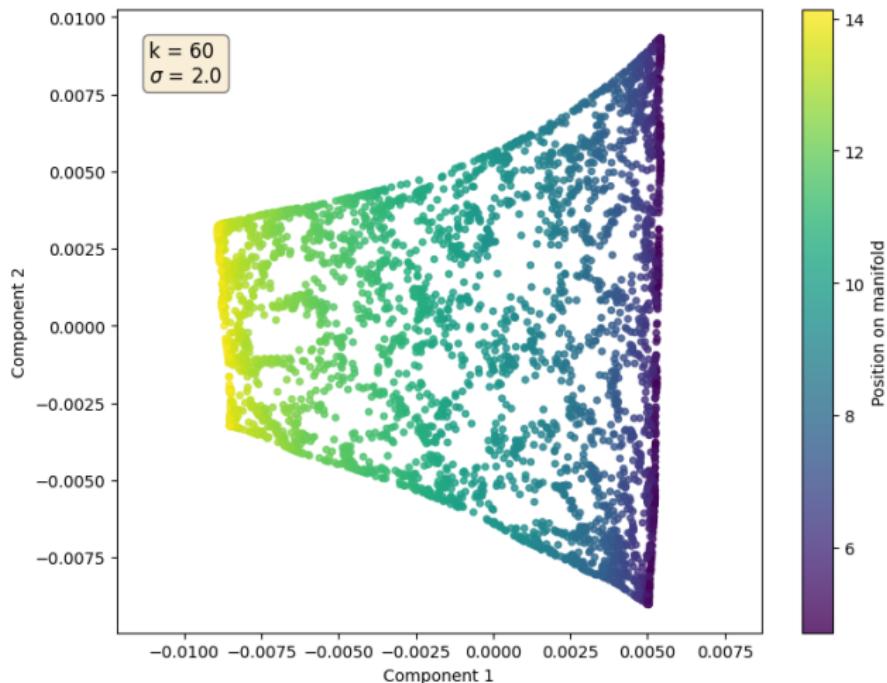


Figure: 2D Embedding of the Swiss Roll Dataset using Laplacian Eigenmaps

Effect of Neighbors (k) on Embedding

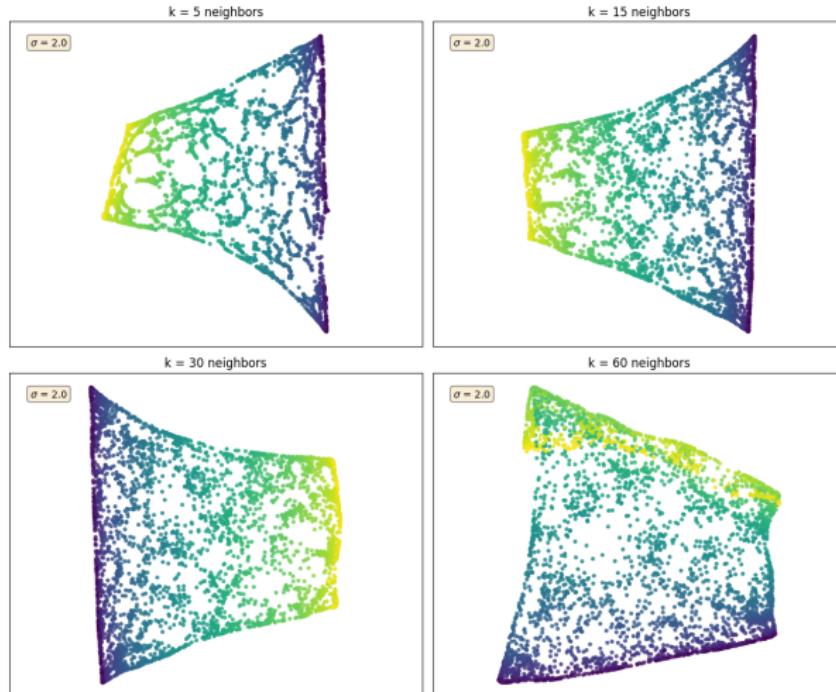


Figure: Effect of Varying Number of Neighbors (k) on the Embedding

Effect of Sigma (Kernel Bandwidth) on Embedding

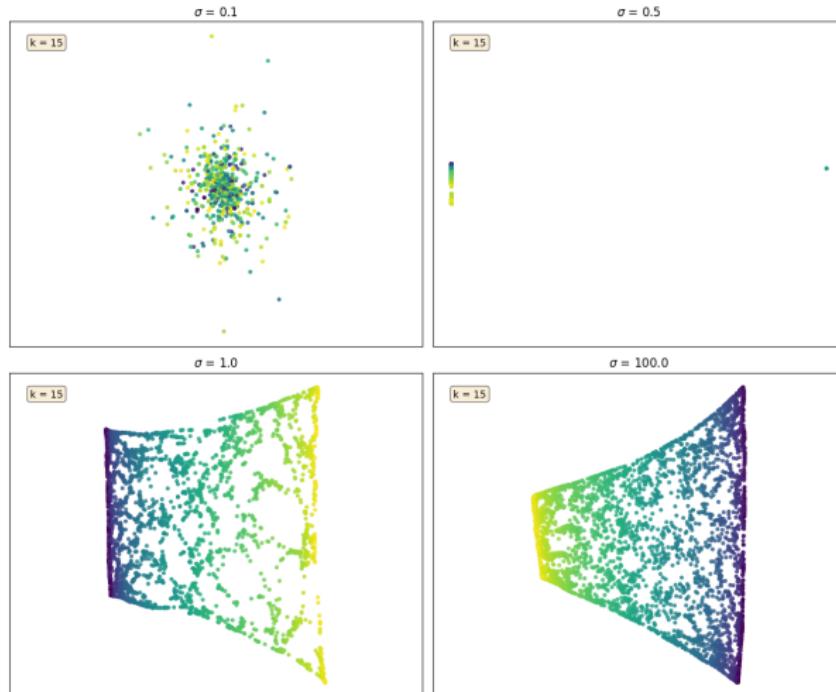


Figure: Effect of Varying Sigma (Kernel Bandwidth) on the Embedding

PCA vs Laplacian Eigenmaps on Swiss Roll

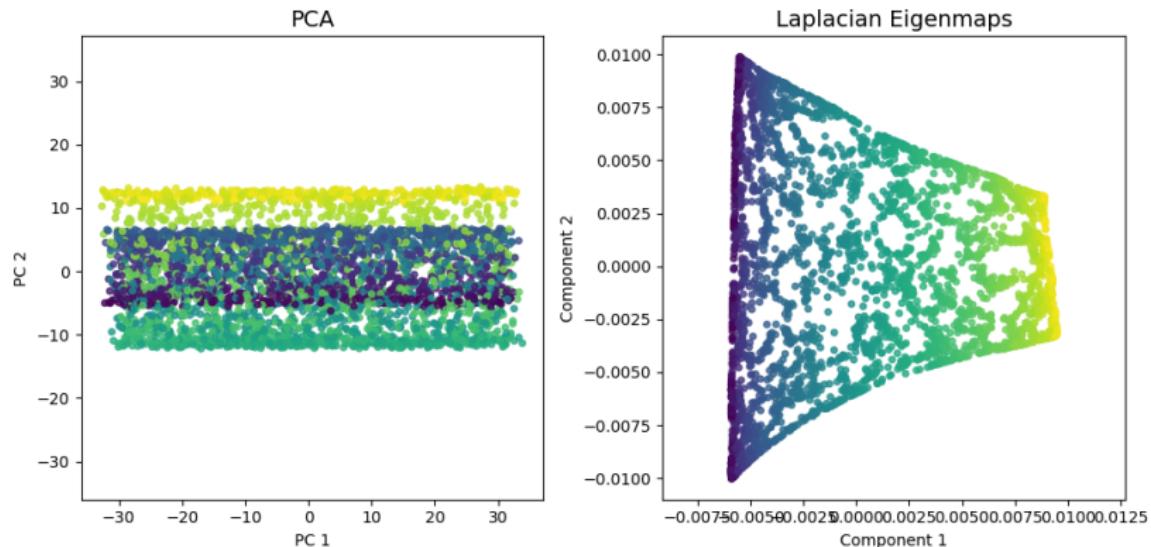


Figure: Comparison of PCA and Laplacian Eigenmaps on the Swiss Roll Dataset

Mammoth Dataset

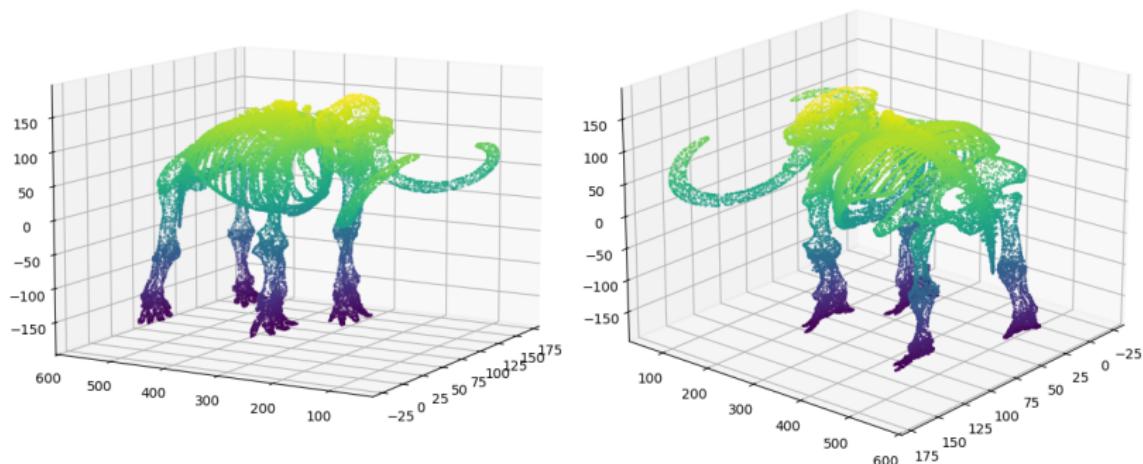


Figure: 3D Mammoth Dataset Visualization

Laplacian Eigenmaps Embedding of Mammoth Dataset

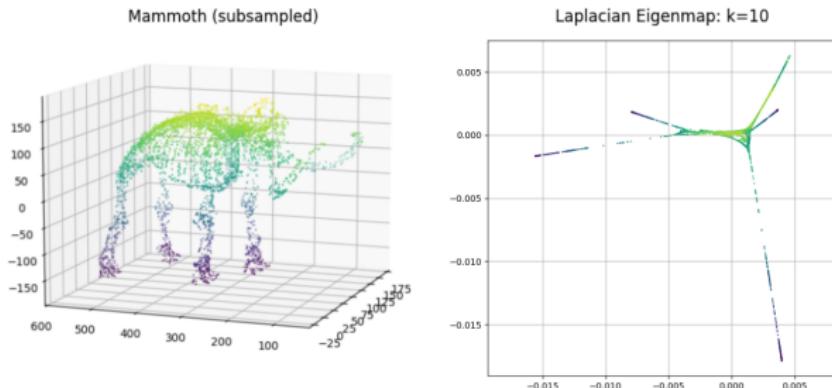


Figure: Laplacian Eigenmaps Embedding of the Mammoth Dataset

Observation

The effect shown on the picture is called the starfish effect. It reflects perfectly how the Laplacian Eigenmaps preserves connectivity, but it fails to preserve global structure.

Laplacian Eigenmaps as UMAP initialization

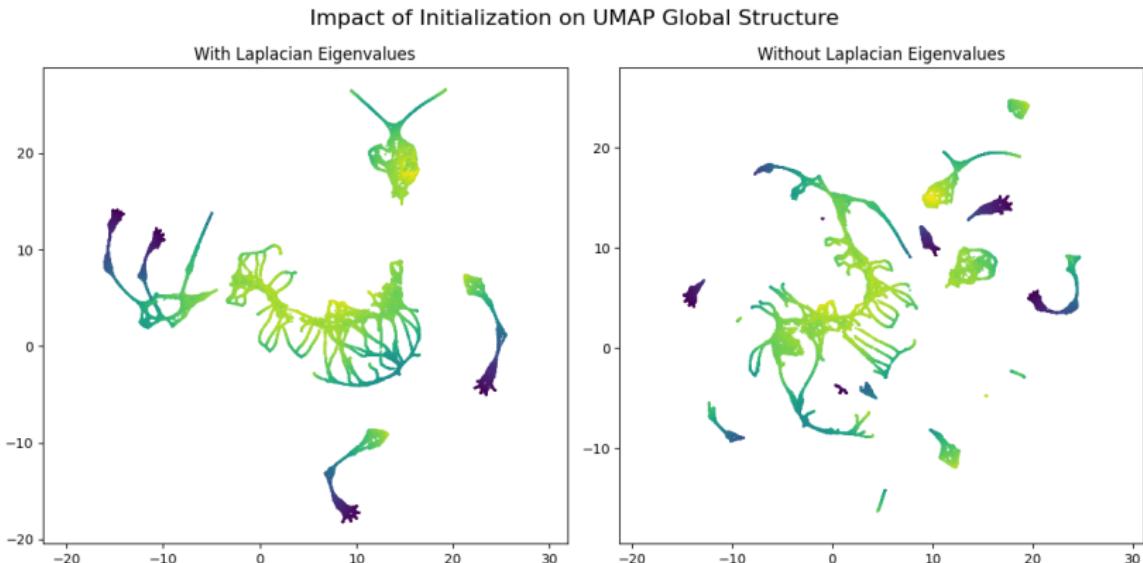


Figure: UMAP Embedding of the Mammoth Dataset with Laplacian Eigenmaps Initialization and Random Initialization