

Proyecto 7: Laplacian Eigenmaps para reducción de dimensión y embeddings

Asignatura: Métodos Diferenciales para la IA

Sergio Herreros Pérez, Mario Kroll Merino, Carlos Mazuecos Reíllo,
José Andrés Ridruejo Tuñón

Why Laplacian Eigenmaps?

- Classical dimensionality reduction techniques (PCA, MDS) are **linear** and fail to capture non-linear structures.
- Real datasets often lie near a **non-linear manifold** embedded in high dimensions.
- **Laplacian Eigenmaps** aim to preserve the **local geometry** of the data.

Key Idea

Build a graph that captures local relationships and use the eigenvectors of the **graph Laplacian** to obtain a smooth low-dimensional embedding.

Laplacian Eigenmaps: Core Idea

- ① Build a **neighborhood graph** using k-NN or a kernel.
- ② Assign similarity weights w_{ij} between connected points.
- ③ Construct the graph Laplacian: $L = D - A$.
- ④ Solve the spectral problem:

$$Lf = \lambda f.$$

Final Embedding

The eigenvectors associated with the smallest non-zero eigenvalues provide the embedding coordinates.

Geometric Interpretation

- The graph Laplacian measures the **smoothness** of a function on the graph.
- Minimizing $f^\top L f = \sum_{i,j} w_{ij}(f_i - f_j)^2$ encourages nearby nodes to have similar values.
- The smallest-energy eigenvectors correspond to the "smoothest" functions on the graph.

Why this reduces dimension

These smooth eigenfunctions capture the local geometric structure of the underlying manifold.

The Continuous Problem: Laplace–Beltrami

On a continuous domain or manifold:

$$\min_{u: \|u\|=1} \int_{\Omega} |\nabla u|^2 dx$$

leads to the eigenvalue problem:

$$-\Delta u = \lambda u.$$

- The eigenvalues encode the **geometry** of the manifold.
- The eigenfunctions are smooth modes with minimal oscillation.

From Continuous to Discrete: The Key Connection

Continuous Laplacian

$$-\Delta u = \lambda u$$

$$\int |\nabla u|^2$$

Graph Laplacian

$$Lf = \lambda f$$

$$f^\top L f = \sum w_{ij} (f_i - f_j)^2$$

Conceptual Correspondence

- $|\nabla u|^2 \longleftrightarrow (f_i - f_j)^2$
- Laplace–Beltrami operator \longleftrightarrow Graph Laplacian
- Smoothness on the manifold \longleftrightarrow Smoothness on the graph

Laplacian Eigenmaps in Graph-Based Learning

- Laplacian Eigenmaps provide the spectral foundation for many graph learning algorithms:
 - **Spectral Clustering:** uses the smallest non-zero Laplacian eigenvectors to partition graphs.
 - **Graph Convolutional Networks (GCNs):** graph convolutions are defined using the Laplacian spectrum.
- The Laplacian encodes the local geometry of the dataset by capturing similarity relations through weights w_{ij} .
- Eigenvectors provide smooth representations that respect this geometry.

Connection to the Finite Element Method (FEM)

Shared Variational Principle

Both FEM and Laplacian Eigenmaps are based on minimizing an energy associated with the Laplacian:

$$\int_{\Omega} |\nabla u|^2 \longleftrightarrow \sum_{i,j} w_{ij}(f_i - f_j)^2 = f^T L f.$$

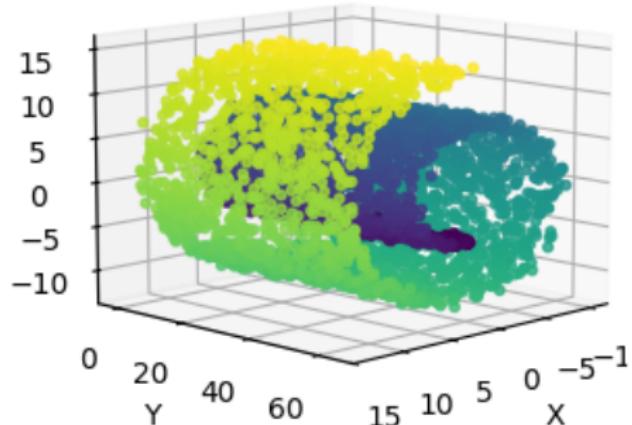
- FEM approximates Δ on a mesh; Laplacian Eigenmaps approximate Δ on a **similarity graph**.
- In both cases, low-energy eigenfunctions correspond to the smoothest modes.
- This provides a rigorous bridge between PDE-based models and graph learning.

From Continuous Geometry to Discrete Data

- Laplacian Eigenmaps allow transferring ideas from differential geometry to data analysis:
 - **Diffusion:** heat flow on a manifold \leftrightarrow diffusion processes on graphs.
 - **Smoothness:** low-oscillation eigenfunctions \leftrightarrow low-variation graph signals.
 - **Eigenmodes:** Laplace–Beltrami eigenfunctions \leftrightarrow graph Laplacian eigenvectors.
- This creates a unified framework connecting:
 - PDEs and variational principles,
 - manifold geometry,
 - and machine learning on graphs.
- The result: geometric structure of data becomes accessible even in discrete, high-dimensional settings.

Swiss Roll Dataset

Swiss Roll - View 1



Swiss Roll - View 2




Figure: 3D Visualization of the Swiss Roll Dataset

Laplacian Eigenmaps Embedding of Swiss Roll

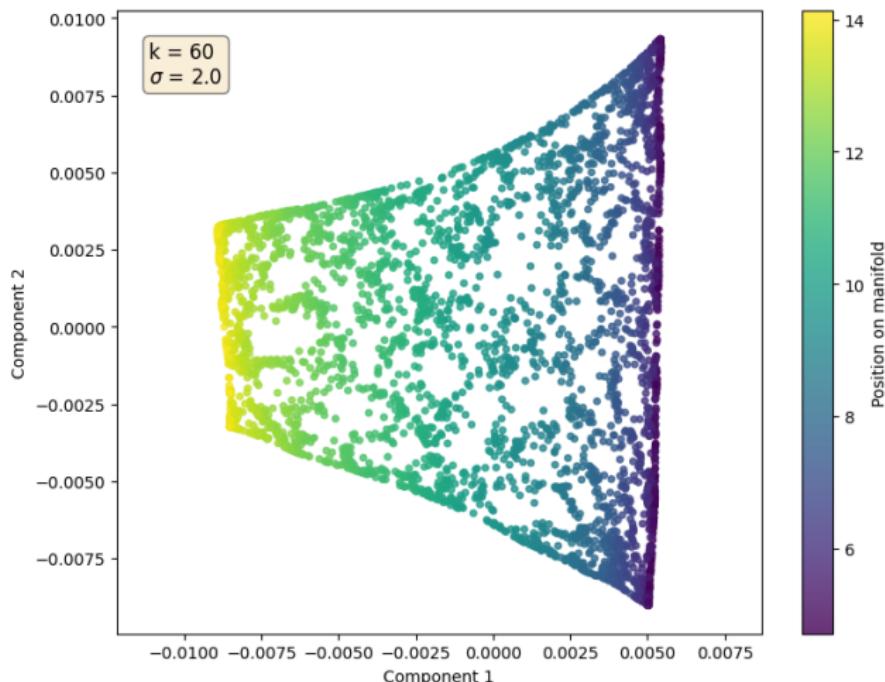


Figure: 2D Embedding of the Swiss Roll Dataset using Laplacian Eigenmaps

Effect of Neighbors (k) on Embedding




Figure: Effect of Varying Number of Neighbors (k) on the Embedding

Effect of Sigma (Kernel Bandwidth) on Embedding

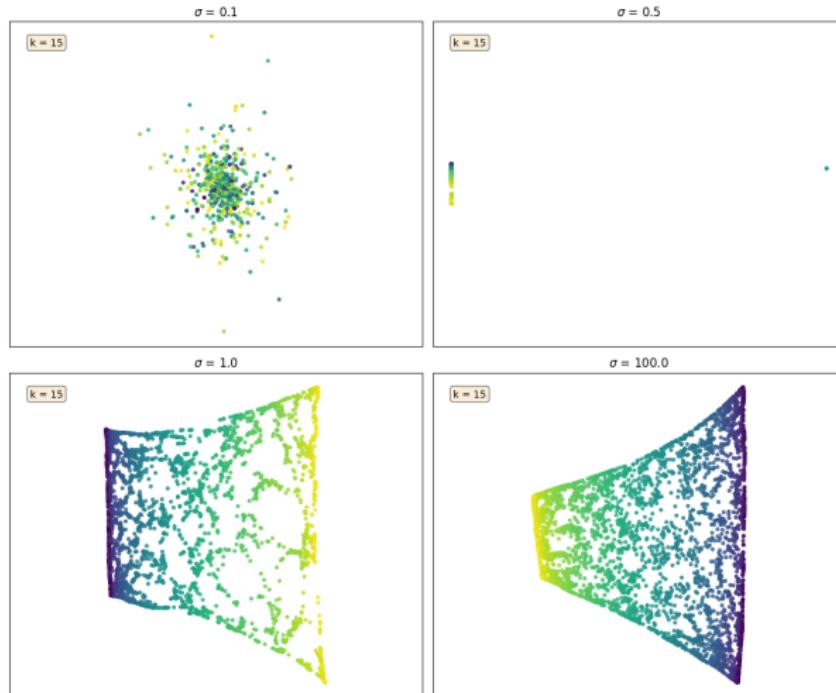


Figure: Effect of Varying Sigma (Kernel Bandwidth) on the Embedding

PCA vs Laplacian Eigenmaps on Swiss Roll

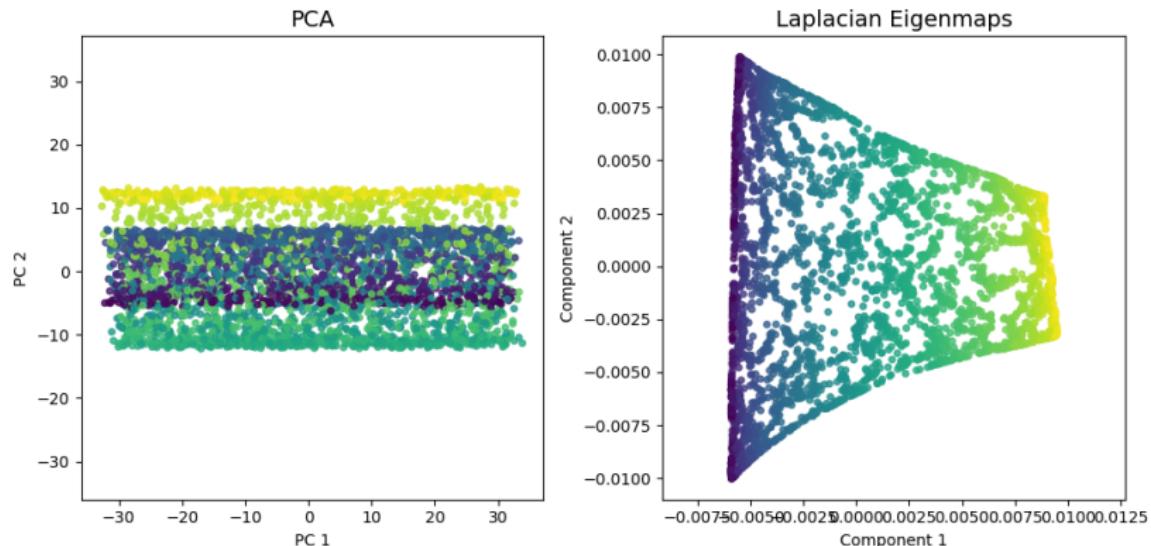


Figure: Comparison of PCA and Laplacian Eigenmaps on the Swiss Roll Dataset

Mammoth Dataset




Figure: 3D Mammoth Dataset Visualization

Laplacian Eigenmaps Embedding of Mammoth Dataset




Figure: Laplacian Eigenmaps Embedding of the Mammoth Dataset

Observation

The effect shown on the picture is called the starfish effect. It reflects perfectly how the Laplacian Eigenmaps preserves connectivity, but it fails to preserve global structure.

Laplacian Eigenmaps as UMAP initialization




Figure: UMAP Embedding of the Mammoth Dataset with Laplacian Eigenmaps Initialization and Random Initialization